

TONE

Promoting empathy among Twitter Users, in order to reduce offensive content that harms the wellness of others.

Our Team



Noor Gill

Project Manager
SME & EDA Lead



Evan Phillips

Machine Learning Lead
Technical Product Manager



Alejandro Pelcastre

Lead Data Engineer
Co-Web Developer



Jaypal Bhatia

Product Manager
Data Analyst



Elias Saravia

Web Developer
UI Designer

Presentation Overview



The Problem

1

Product Demo

2

Technical Approach

3

Potential Next Steps

4

The Problem



Increasing Personal attacks

Created an environment on social media where others are targeted for their race, sexual orientation, disabilities, and much more.

Twitter tries to identify offensive content

but struggles to distinguish between the types of personal attacks, experimented with features

The Problem in Numbers



3.8 Million

Tweets removed for offensive content policy
July 2020-August 2020


1.1 Million

Twitter users disciplined for offensive content



Eric

Eric is an impulsive tweeter, who often uses derogatory sexist language casually. Even though his intention may be joking, he often replies to friends like this.



Emily

Emily is an online friend of Eric, who is impacted by Eric's hurtful tweets. Eric often replies to Emily, and this has impacted her mental health negatively.



Sample Dataset Tweet

@username



Rudy Gay a fa***t for not changing his name once he got money

12:00 PM · Jun 1, 2021

LGBTQ+ Phobic



Sample Dataset Tweet

@username



Twitter game is on point tonight btw guys, if you haven't caught on you're full r*****d

12:00 PM · Jun 1, 2021

Disability Shaming



Sample Dataset Tweet

@username



Every time I try to quit smoking, some dumb b***h always gotta be f*****g annoying and try my patience lmao

12:00 PM · Jun 1, 2021

Sexist



Sample Dataset Tweet

@username



The thing about working with these black people. These b*****s never on f*****g time then wanna complain about a short check. B**** die. today. now.

12:00 PM · Jun 1, 2021

Racist, sexist

Benefits to Solving This Problem

Less Attacks. More Empathy


Promote more empathy among twitter users by providing detailed feedback of their tweets

Stop Targeting to Stop Declining Mental Health

Reduce targeting online, which is tied to issues such as suicide and declining mental health

Expand Twitter's Current Stance

Expand Twitter's current stance on offensive content to be more comprehensive



Benefits to Solving This Problem

Less Attacks. More Empathy


Promote more empathy among twitter users by providing detailed feedback of their tweets

Stop Targeting to Stop Declining Mental Health

Reduce targeting online, which is tied to issues such as suicide and declining mental health

Expand Twitter's Current Stance

Expand Twitter's current stance on offensive content to be more comprehensive



Benefits to Solving This Problem

Less Attacks. More Empathy


Promote more empathy among twitter users by providing detailed feedback of their tweets

Stop Targeting to Stop Declining Mental Health

Reduce targeting online, which is tied to issues such as suicide and declining mental health

Expand Twitter's Current Stance

Expand Twitter's current stance on offensive content to be more comprehensive



Twitter's Current Stance

“Healthy conversation is a shared responsibility. If your Tweet reply is identified as using potentially harmful or offensive language, we may ask you, via a prompt, if you want to review it before sending.”



A More Expansive Stance

Twitter and Tone work together to identify the specific type of harmful content you may be tweeting. You'll receive a breakdown of your tweet and the categories of offensive content present in your tweet.



Current and Future Impact

Social Media Beyond Twitter

Market size is not limited to just Twitter, but other social media outlets where this problem is prevalent.

Negative Tweets Still Growing

July 2020-August 2020, 3.8 million tweets removed for Twitter's offensive content policy.

Beta Feature to Edit Tweets

Twitter's current venture to edit offensive tweets shows movement and need for expansion regarding offensive content



Current and Future Impact

Social Media Beyond Twitter

Market size is not limited to just Twitter, but other social media outlets where this problem is prevalent.

Negative Tweets Still Growing

July 2020-August 2020, 3.8 million tweets removed for Twitter's offensive content policy.

Beta Feature to Edit Tweets

Twitter's current venture to edit offensive tweets shows movement and need for expansion regarding offensive content



Current and Future Impact

Social Media Beyond Twitter

Market size is not limited to just Twitter, but other social media outlets where this problem is prevalent.

Negative Tweets Still Growing

July 2020-August 2020, 3.8 million tweets removed for Twitter's offensive content policy.

Beta Feature to Edit Tweets

Twitter's current venture to edit offensive tweets shows movement and need for expansion regarding offensive content



Product Demo

Offensive Content Warning

TONE

Toning Down Bad Vibes

Input your tweet below:

Demo Analysis



Name

@Username

...

@SwaaggyyV: Fucking chinks in Clash of Clans always fuck me up. Bitch ass kamikaze faggots.

12:00 PM · Jun 1, 2021










Tone Representation



Our Product



Create a platform that identifies targeting tweets in the following categories as a proof of concept:

Neutral



General Criticism



Disability Discrimination



Racial Prejudice



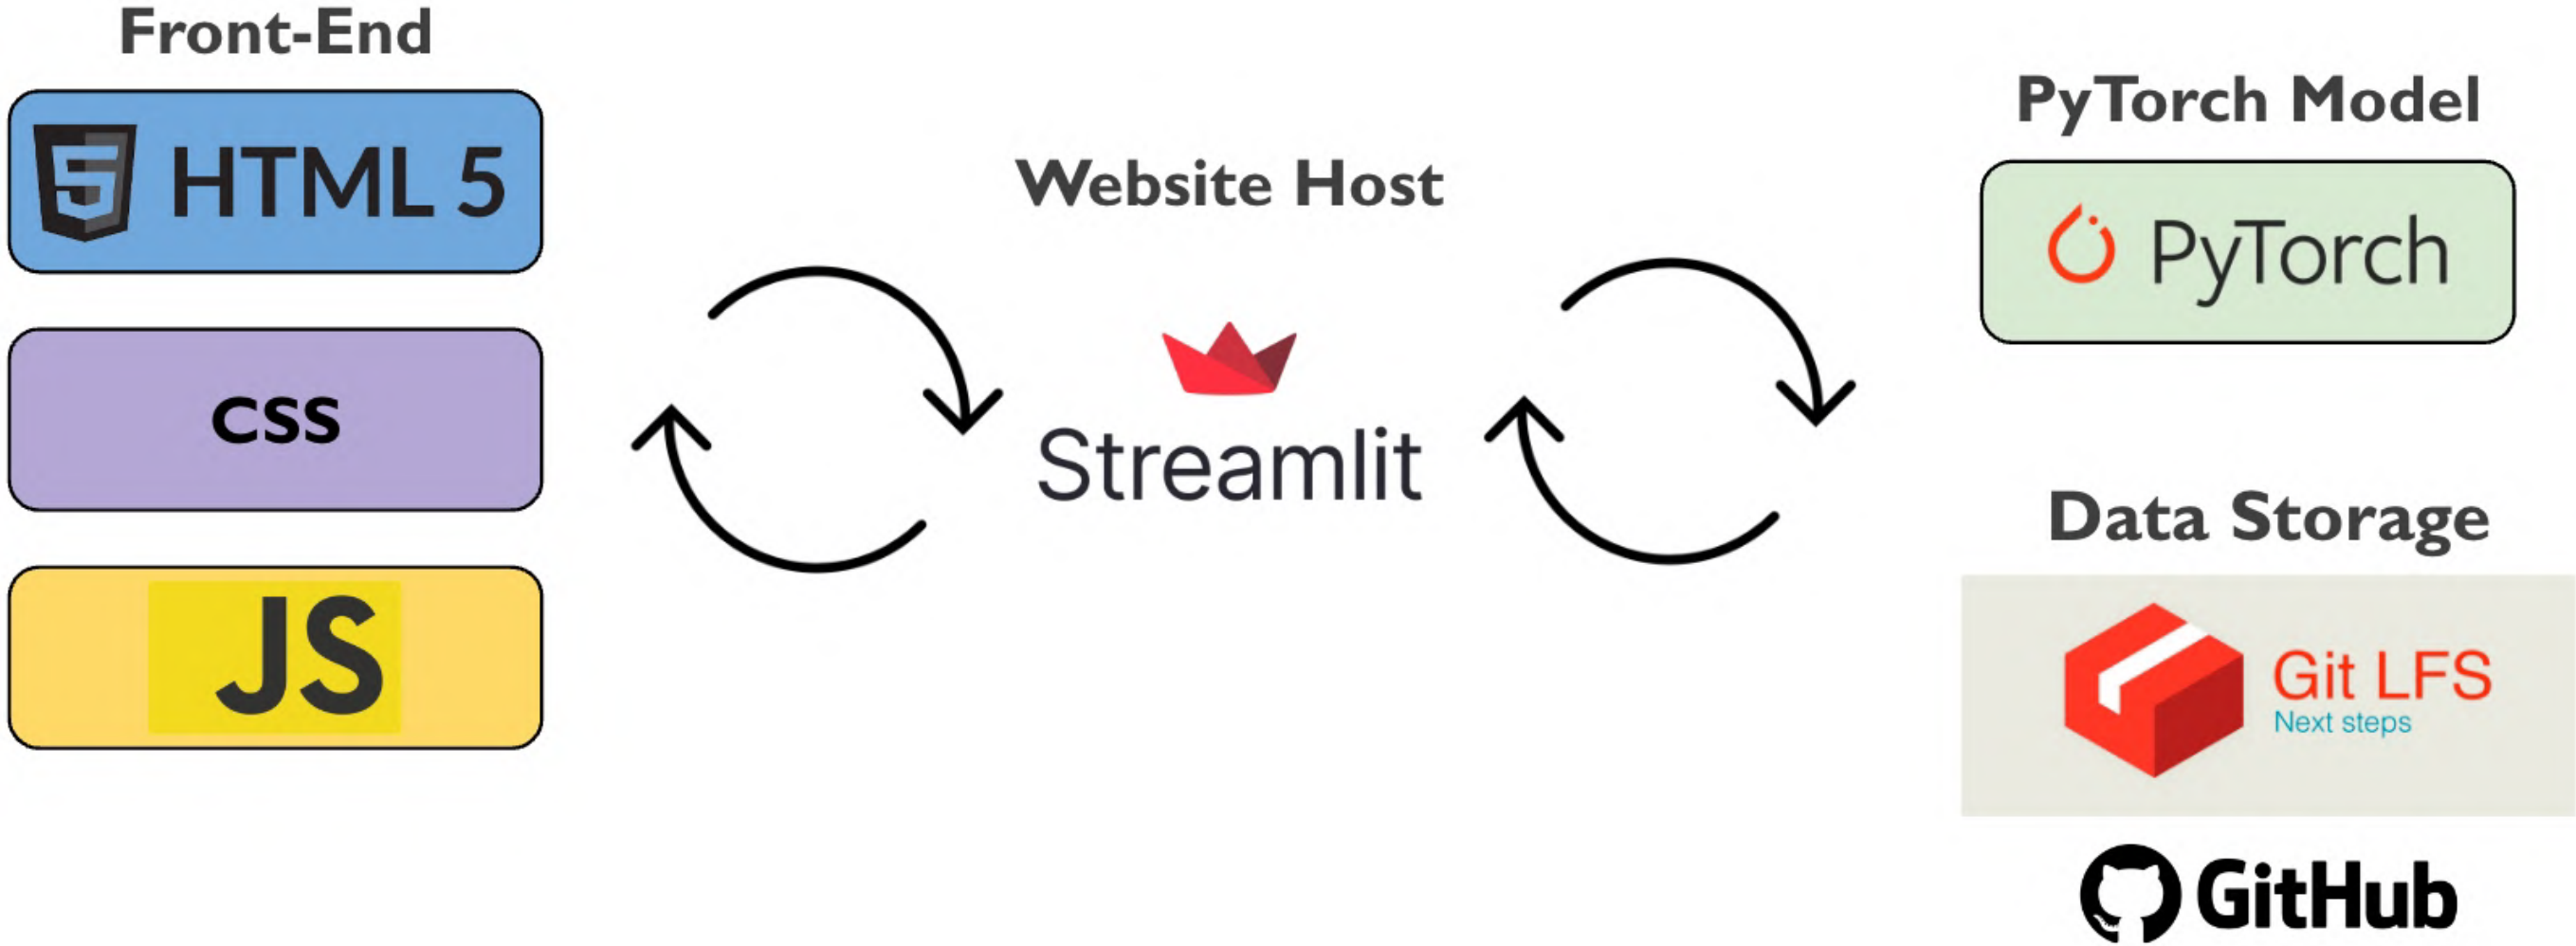
Sexism



LGBTQ+ Phobia



Data Engineering Pipeline



Data Labeling



Sample Dataset Tweet
@username

The thing about working with these black people. These b*****s never on f*****g time then wanna complain about a short check. B***** die. today. now.

12:00 PM · Jun 1, 2021

- (0) Neutral
- (1) General Criticism
- (2) Disability Shaming
- (3) Racial Prejudice
- (4) Sexism
- (5) LGBTQ+ Phobia

Tweet	(0)	(1)	(2)	(3)	(4)	(5)
Person A	0	1	0	1	0	0
Person B	0	1	0	0	0	0
Person C	0	1	0	1	0	0
Person D	0	0	0	1	0	0
Person E	0	1	0	1	0	0

Pre-Processing Data

Removal of duplicate tweets and NA values

- Deleted 836 tweets in total

Removal of URLs, hashtags, usernames, emojis, numbers, and RT

- https://~
- #...
- @...
- 1234...
- RT...
- 🤔

Implemented a 90-5-5 train-validation-test split of our dataset

- Train size: 20,372 tweets
- Validation size: 1,132 tweets
- Test size: 1,132 tweets


Tokenization (cased-BERT tokenizer)

- **ex:** ['[CLS]', 'For', 'the', 'record', 'No', '##H', '##omo', 'but', 'don', '##t', 'care', 'who', 'is', 'unless', 'I', 'gotta,']
- "bad" vs. "BAD"

Base Framework

Key Takeaways:

- Train for 4 epochs
- Batch = 32
- Learning rate = $2e-5 = 0.00002$
- Maximum token length = 50





Token Length Distribution

Dynamism of our Model

PyTorch Lightning:

- ModelCheckpoint
- Simulate 100 training steps
- Optimal number of 2 epochs




Learning Rate vs. Step per Epoch

Model Evaluation

Key Takeaways:

- 98.4% Accuracy
- Sigmoid, not Softmax Output Layer
- ROC fit depends on class balance




	precision	recall	f1-score	support
disability shaming	1.00	0.84	0.91	19
racial prejudice	0.99	0.99	0.99	97
sexism	1.00	1.00	1.00	747
lgbtq+ phobia	1.00	1.00	1.00	66
micro avg	1.00	0.99	1.00	929
macro avg	1.00	0.96	0.98	929
weighted avg	1.00	0.99	1.00	929
samples avg	0.72	0.71	0.72	929

Model Limitations



Tone Representation



What's Next for **TONE**?

Increase Accuracy by 15%

Optimization + Collect More Training Data

Improve Front End Development

New Host + Improve User Experience

Optimize Data Pipeline Infrastructure

Leverage AWS for data streaming + storage

Conduct Usability Testing + Feedback

Collect user interviews, surveys, and focus groups

Implement New Highlight Feature

Highlighting / Detecting Hate Speech Words

Research Contextualization and Labeling

Communicating with experts to improve product

5 Key Takeaways

Product Must Reflect the Source of the Data



NLP + ML Model Development, Bias, and Optimization



Conducting Research + Talking to Experts is Important



How to Create a Startup + Its Components



The 4 Cs: Consistent Communication, Contribution, and Collaboration



WE ARE T NE

Our Mission:

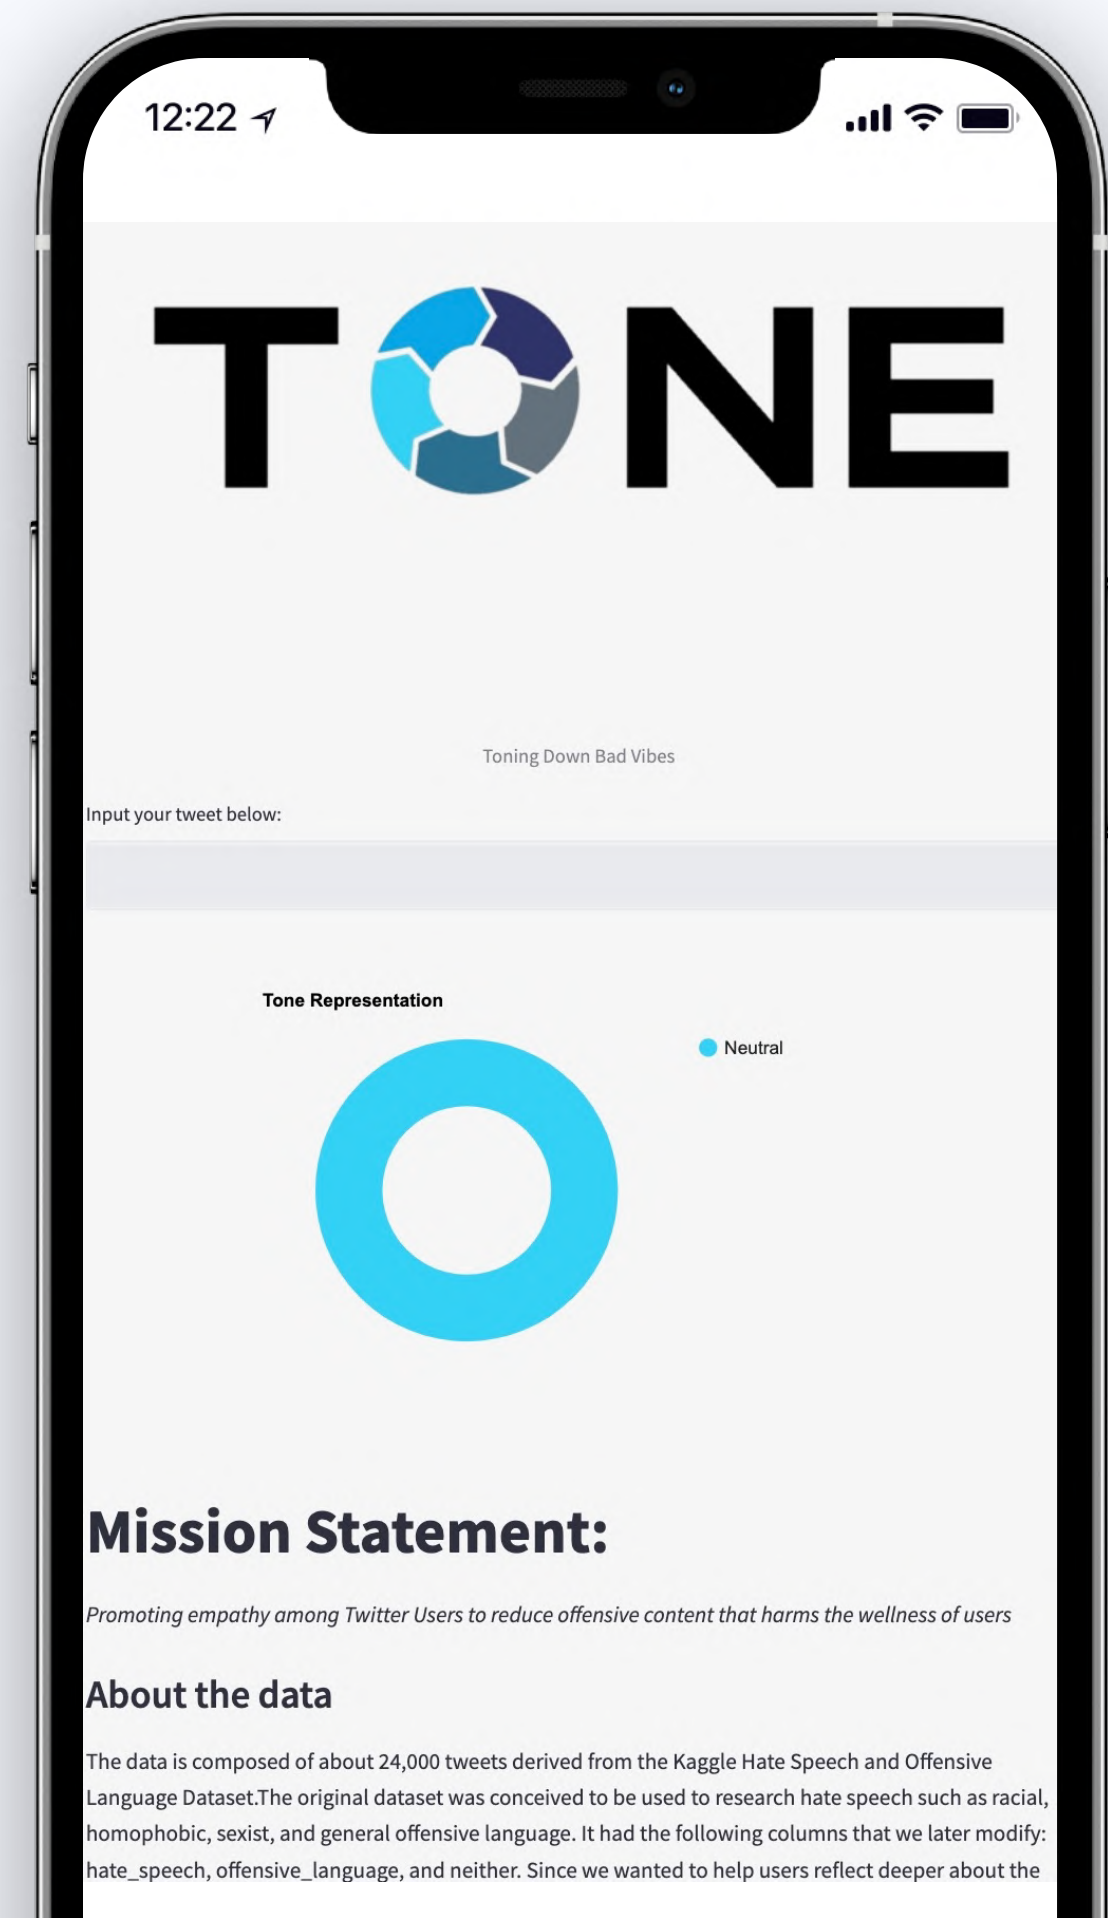
To promote empathy among Twitter Users, in order to reduce offensive content that harms the wellness of others.

Company Name : Date
Acme IncAugust 2024

Contact:
email@acme.com

Thank You!

Any Questions?



Acknowledgements

We would like to give special thanks to the following individuals for helping us out with our development:

- Prof. Joyce Shen
- Prof. Zona Kostic
- Prabhu Narsina
- Kevin Hartman
- Robert Wang (AWS)
- UC Berkeley 5th Year MIDS Cohort 2022